

## Homework #5

- a) **(5 points)** Modify the chained hash table from the book (`chtbl.h` and `chtbl.c`) so that it auto-grows when its load factor exceeds a given value. (Auto-growing means increasing the number of buckets in the hash table then rehashing the existing elements.)

Begin by modifying the `chtbl_init` function from the book to have the following prototype:

```
int chtbl_init(CHTbl *htbl,
              int buckets,
              int (*h)(const void *key),
              int (*match)(const void *key1, const void *key2),
              void (*destroy)(void *data),
              double maxLoadFactor,
              double resizeMultiplier);
```

The `htbl`, `buckets`, `h`, `match`, and `destroy` parameters have the same meaning as defined in the book. The `maxLoadFactor` parameter is the maximum load factor the hash table should be allowed to reach before being auto-resized. The `resizeMultiplier` parameter is the amount by which the number of buckets should be multiplied when a resize occurs (this value must be greater than 1). The `maxLoadFactor` and `resizeMultiplier` values should be stored in new fields in the `CHTbl` struct.

Next, modify the code in the `chtbl_insert` function as follows:

1. Resize the hash table when the load factor exceeds the maximum load factor. The new size of the hash table should be the old size times the `resizeMultiplier`. All elements currently in the hash table must be rehashed and placed into new buckets.
2. Change the method by which hash codes are mapped to buckets to use the multiplication method instead of the division method.

Modify any other parts of `CHTbl` as needed. Note that removing items from the hash table should not cause the number of buckets in the table to shrink.

- b) **(3 points)** Implement a program that demonstrates inserts and lookups with an auto-resizing hash table. This program should initialize the hash table to a small number of buckets then begin inserting integers until a resize occurs. After each insert the program should output the following information:

- Number of buckets in the table
- Number of elements in the table
- The table's load factor
- The table's max load factor
- The table's resize multiplier

For example, if the hash table was initialized to start with 5 buckets, a max load factor of 0.5, and a resize multiplier of 2.0, the following output should be displayed as elements are inserted:

```
buckets 5, elements 1, lf 0.20, max lf 0.5, resize multiplier 2.0
buckets 5, elements 2, lf 0.40, max lf 0.5, resize multiplier 2.0
buckets 10, elements 3, lf 0.33, max lf 0.5, resize multiplier 2.0
```

Note that in the 3<sup>rd</sup> line of output the number of buckets has doubled. This happened because the load factor that would result from inserting the 3<sup>rd</sup> element would have caused the load factor to exceed the max load factor (0.5) so the hash table was auto-resized.

After the resize occurs, your program must demonstrate successfully looking up a value that was inserted before the resize and must also demonstrate unsuccessfully looking up a value that does not exist in the table.

c) **(1 point)** Answer the following questions:

1. What is the Big-O execution performance of an insert now that auto-resizing can take place?
2. Why do you think you were required to change `chtbl_insert` to use the multiplication method instead of the division method to map hash codes to buckets?

d) **(1 point)** Make sure your source code is well-commented, consistently formatted, uses no magic numbers/values, follows programming best-practices, and is ANSI-compliant.

**Turn in all source code, program output, diagrams, and answers to questions in a single Word or PDF document.**